

Cifar-100 实验总结

训练展示:

| 测试代码 | 网络结构 | nb_epoch | loss | acc | val_loss | val_acc | 训练时间/s |
|------------------|--|----------|--------|--------|----------|---------|--------|
| cifar100_cnn_dl | 参考结构 | 50 | 2.4899 | 0.3609 | 2.3722 | 0.3856 | 1641 |
| cifar100_cnn_dl1 | 卷积核: 3*3->5*5 最后一个激活函数: 'relu'->'tanh' SGD: lr = 0.01 -> lr = 0.001 增加卷积层: 第二部分 | 50 | 1.8501 | 0.4938 | 1.7869 | 0.5177 | 2646 |

cifar100_cnn_dl 网络结构:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3, border_mode='same',
                        input_shape=(img_channels, img_rows, img_cols)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 3, 3, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

# let's train the model using SGD + momentum (how original).
sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd)
```

cifar100_cnn_dl 结果:

```
Epoch 45/50
50000/50000 [=====] - 33s - loss: 2.5053 - acc: 0.3545 - val_loss: 2.2624 - val_acc: 0.4114
Epoch 46/50
50000/50000 [=====] - 33s - loss: 2.5246 - acc: 0.3518 - val_loss: 2.2108 - val_acc: 0.4231
Epoch 47/50
50000/50000 [=====] - 33s - loss: 2.5068 - acc: 0.3538 - val_loss: 2.2649 - val_acc: 0.4095
Epoch 48/50
50000/50000 [=====] - 33s - loss: 2.5038 - acc: 0.3574 - val_loss: 2.2118 - val_acc: 0.4264
Epoch 49/50
50000/50000 [=====] - 33s - loss: 2.5000 - acc: 0.3548 - val_loss: 2.2677 - val_acc: 0.4133
Epoch 50/50
50000/50000 [=====] - 33s - loss: 2.4899 - acc: 0.3609 - val_loss: 2.3722 - val_acc: 0.3856
root@young-pc: /home/young/ImageClassify#
```

cifar100_cnn_dl1 网络结构:

```
model = Sequential()
model.add(Convolution2D(32, 5, 5, border_mode='same',
                        input_shape=(img_channels, img_rows, img_cols)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 5, 5))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 5, 5))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('tanh'))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

# let's train the model using SGD + momentum (how original).
sgd = SGD(lr=0.001, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd)
```

cifar100_cnn_dl1 结果:

```
Epoch 45/50
50000/50000 [=====] - 54s - loss: 1.9008 - acc: 0.4793 - val_loss: 1.8743 - val_acc: 0.4967
Epoch 46/50
50000/50000 [=====] - 54s - loss: 1.8933 - acc: 0.4827 - val_loss: 1.8440 - val_acc: 0.5047
Epoch 47/50
50000/50000 [=====] - 54s - loss: 1.8754 - acc: 0.4848 - val_loss: 1.8288 - val_acc: 0.5069
Epoch 48/50
50000/50000 [=====] - 54s - loss: 1.8687 - acc: 0.4862 - val_loss: 1.8331 - val_acc: 0.5066
Epoch 49/50
50000/50000 [=====] - 54s - loss: 1.8624 - acc: 0.4886 - val_loss: 1.8274 - val_acc: 0.5077
Epoch 50/50
50000/50000 [=====] - 54s - loss: 1.8501 - acc: 0.4938 - val_loss: 1.7869 - val_acc: 0.5177
root@young-pc: /home/young/ImageClassify#
```

根据 cifar10 得到的最好测试模型训练 cifar100 数据集得到的效果, 比根据 cifar10 参考模型训练 cifar100 的效果要好很多, 提高了 13 个百分点, 但是, 两者的不足是测试成功率和交叉验证成功率都较低。如果需要对 cifar100 数据集训练和测试, 还需要探索新的深度学习网络结构。

测试展示:

| 测试代码 | 网络结构 | 测试数据 | loss | acc | 测试时间/s |
|------------------|--|-------------------------|--------|--------|--------|
| cifar100_cnn_dl | 参考结构 | Cifar100 TestSet: 10000 | 2.3722 | 0.3856 | 1 |
| cifar100_cnn_dl1 | 卷积核: 3*3->5*5 最后一个激活函数: 'relu'->'tanh' SGD: lr = 0.01 -> lr = 0.001 增加卷积层: 第二部分 | Cifar100 TestSet: 10000 | 1.7869 | 0.5177 | 2 |

cifar100_cnn_dl 网络结构:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3, border_mode='same',
                        input_shape=(img_channels, img_rows, img_cols)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 3, 3, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

# let's train the model using SGD + momentum (how original).
sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd)
```

cifar100_cnn_dl 测试结果:

```
root@young-pc:/home/young/ImageClassify# THEANO_FLAGS=mode=FAST_RUN,device=cpu,floatX=float32 python cifar100_cnn_dl_predict.py
Using Theano backend.
Using gpu device 0: GeForce GTX 960M (CNMeM is disabled, CuDNN 3007)
X_test shape: (10000, 3, 32, 32)
10000 test samples
10000/10000 [=====] - 1s
test score: - loss: 2.3722 - acc: 0.3856
10000/10000 [=====] - 1s
test predict_classes: [79 80 45 ... 96 42 70]
root@young-pc:/home/young/ImageClassify#
```

cifar100_cnn_dl1 网络结构:

```
model = Sequential()
model.add(Convolution2D(32, 5, 5, border_mode='same',
                        input_shape=(img_channels, img_rows, img_cols)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 5, 5))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 5, 5))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(Convolution2D(64, 5, 5, border_mode='same'))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('tanh'))
model.add(Dense(nb_classes))
model.add(Activation('softmax'))

# let's train the model using SGD + momentum (how original).
sgd = SGD(lr=0.001, decay=1e-6, momentum=0.9, nesterov=True)
model.compile(loss='categorical_crossentropy', optimizer=sgd)
```

cifar100_cnn_dl1 测试结果:

```
root@young-pc:/home/young/ImageClassify# THEANO_FLAGS=mode=FAST_RUN,device=cpu,floatX=float32 python cifar100_cnn_dl1_predict.py
Using Theano backend.
Using gpu device 0: GeForce GTX 960M (CNMeM is disabled, CuDNN 3007)
X_test shape: (10000, 3, 32, 32)
10000 test samples
10000/10000 [=====] - 2s
test score: - loss: 1.7869 - acc: 0.5177
10000/10000 [=====] - 2s
test predict_classes: [30 80 55 ... 51 42 70]
root@young-pc:/home/young/ImageClassify#
```

小结: 实验表明, 改进过的结构可以将测试成功率提高 13 个百分点, 效果还是不错的。

附录:

1. 文档中所有源文件都在文件夹 ImageClassifySummary_Cifar100 里, 实验结果截图在 resultimage 文件夹下。
2. 保存的模型和测试结果均以其对应源代码文件名区别。例如, 源代码为 cifar100_cnn_dl.py, 训练保存的模型文件为 my_model_architecture_cifar100_dl.json, 模型权重为 my_model_weights_cifar100_dl.h5, 测试代码为 cifar100_cnn_dl_predict.py, 测试预测结果为 test_predict_cifar100_dl.txt, 测试分类结果为 test_predict_classes_cifar100_dl.txt。其他源文件的测试结果文件名与之类似。
3. 测试平台参数:
CPU: Inter(R) Core(TM) i7-6700HQ CPU @ 2.60
GPU: GeForce GTX 960M
Cuda: 7.5
CuDNN: 7.0
Python: 2.7.6
Theano: 0.8.1
Keras: 0.3.3